

Benjamin Garcia

(626) 246-4778 | bentgarcia05@gmail.com | [GitHub](#) | [LinkedIn](#) | [Portfolio](#)

EDUCATION

University of California, Los Angeles

Bachelor of Science, Computer Science

Los Angeles, CA

Expected Graduation: Dec. 2027

Relevant Coursework: Data Structures and Algorithms, Operating Systems, Computer Architecture, Programming Languages

TECHNICAL SKILLS

Languages: TypeScript, Python, C++, Java, SQL, JavaScript

Web & UI: React, Next.js, Node.js, GraphQL, WebSockets, React Flow, Tailwind CSS

Backend & Infrastructure: FastAPI, **asyncio**, Temporal, Redis, PostgreSQL, MongoDB, Docker, Linux/Unix

Developer Tools: Git, Graphite, Playwright, Claude Code, Prisma, Zod

EXPERIENCE

Software Engineer Intern

Jun. 2026 – Present

Lindy

San Francisco, CA

- Hardened multi-account Email and Meetings infrastructure serving **27,252 users**, separating OAuth credentials, per-account settings, and agent state to prevent account removal, reconnection, sync, and opt-out flows from silently re-enabling disabled inbox or calendar automation.
- Built account-management and authorization guardrails across web, mobile, and Lindy Teammate for a platform supporting **5,000+ integrations**, covering add, reconnect, delete, rename, and sharing flows while enforcing workspace permissions and secure tool access.
- Improved scheduling and follow-up draft quality across Gmail and Outlook by integrating user instructions, learned writing style, holiday exclusions, account-aware meeting classification, and rendering fixes into production generation flows; helped reduce user-requested scheduling-draft edits from **14.9% to 3.2%** and draft-alert pushback from **6.7% to 3.9%**.
- Shipped Lindy's redesigned home page end to end, reconciling two parallel implementation branches, wiring live meeting, daily-brief, and memory data, adding inbox backprocessing for populated first-use views, and owning Playwright validation, PostHog instrumentation, and staged feature-flag rollout.

Undergraduate Researcher

Dec. 2025 – Present

UCLA Human-Computer Interaction Research Lab

Los Angeles, CA

- Building a **human-in-the-loop** AI research platform that transforms retrieved scientific literature into evidence-grounded agent perspectives, structured debate, and refined research hypotheses.
- Implemented a React Flow hypothesis workspace backed by FastAPI event streams, visualizing **27 workflow event types** across agent reasoning, critique, refinement, and persistence.
- Developed a **5-stage, 20-step** literature-grounded workflow that retrieves Semantic Scholar papers, clusters evidence by embeddings, synthesizes citation-grounded researcher personas, and surfaces provenance in inspectable UI panels.

Software Engineer Intern

Mar. 2025 – Sep. 2025

Todd Agriscience

Los Angeles, CA

- Built production Next.js crop-intelligence dashboards for **5–10 clients**, integrating Palantir Foundry pipelines and AWS Lambda services to deliver real-time agricultural insights.
- Developed data pipelines visualizing **20+ agricultural streams** across soil, irrigation, weather, and yield data to support real-time farm management decisions.

Software Engineer Intern

May 2025 – Jun. 2025

TensorStar

San Francisco, CA

- Engineered real-time React interfaces with Next.js, Redux, and WebSockets, supporting **100+ concurrent users** with **sub-100ms latency**.
- Built secure credential-submission workflows with Python, Redis, and HashiCorp Vault, supporting authentication across **50+ enterprise data sources**.

PROJECTS

PolicyC | TypeScript, Python, asyncio, SQLite, OpenAI API, Zod, FastAPI

Jul. 2026 – Present

- Built a deterministic policy compiler that selects request-specific, dependency-closed subsets from a 43-node, six-domain policy graph, reducing mean model input by up to **98.23%**.
- Engineered a resumable Python **asyncio** runtime that completed three held-out experiments in 37.24 aggregate minutes versus 134.11 minutes of cumulative provider-call latency, achieving **72.2% lower elapsed execution time** and **3.60× effective throughput**.
- Executed **960 GPT-5 mini calls** in paired held-out studies for **\$2.66**, achieving up to **86.49% conditional critical-obligation preservation**; normalized uncached projections estimated an equivalent compiled workload at \$2.43 versus \$6.17 with full prompts—a **60.6% overall cost reduction** and up to **67.8%** in an individual study.